




Making an Internal Kubernetes Offering Generally Available

James Wen - jameswen@spotify.com
@rochesterinnyc

KubeCon + CloudNativeCon North America 2019
November 19th, 2019


About Me

- Senior Site Reliability Engineer
- Once Provisioning + DNS
- Core Infra: Deployment + Kubernetes Migration



Agenda

- Intro Spotify/Ops/Runtime
- GA Journey
- Learnings/Takeaways



Search

MAIN

- Browse
- Activity
- Radio

YOUR MUSIC

- Songs
- Albums
- Artists

PLAYLIST

Discover Weekly

Your weekly mix of discoveries and deep cuts, picked just for you. Updated every Monday, so save your favourites!

Created by: Spotify • 30 songs, 2 hr 04 min

PLAY FOLLOWING

FOLLOWERS

SONG	ARTIST	ALBUM
+ All My Friends	Dylan Kelly	All My Friends
+ Can't Deny My Love	Brandon Flowers	Can't Deny My Love
+ 1000	Ben Khan	1000
+ Fly (FKJ Remix)	June Marleezy	Fly (Remixes) - Single
+ Eventually	Tame Impala	Eventually
+ Abrasive	Ratatat	Abrasive
+ Hear Me Knocking - Alternate Ver...	The Rolling Stones	Can't You Hear Me Knocking
+ ...	Tyler, The Creator	Cherry Bomb
+ ...	Oddisee	That's Love - Single

Discover Weekly

What You Want
Toro Y Moi

1:40

⏮ ⏸ ⏭

🔍 🔊 🔍

Discover Weekly

FOLLOWING

SHUFFLE PLAY

Available Offline

All My Friends
Dylan Kelly • All My Friends

Can't Deny My Love
Brandon Flowers • Can't Deny My Love

What You Want
Toro Y Moi

09:06

Discover Weekly

FOLLOWING

BY SPOTIFY

SHUFFLE PLAY

Available Offline

All My Friends
Dylan Kelly • All My Friends


Can't Deny My Love
Brandon Flowers • Can't Deny My Love

1000
Ben Khan • 1000

What You Want
Toro Y Moi

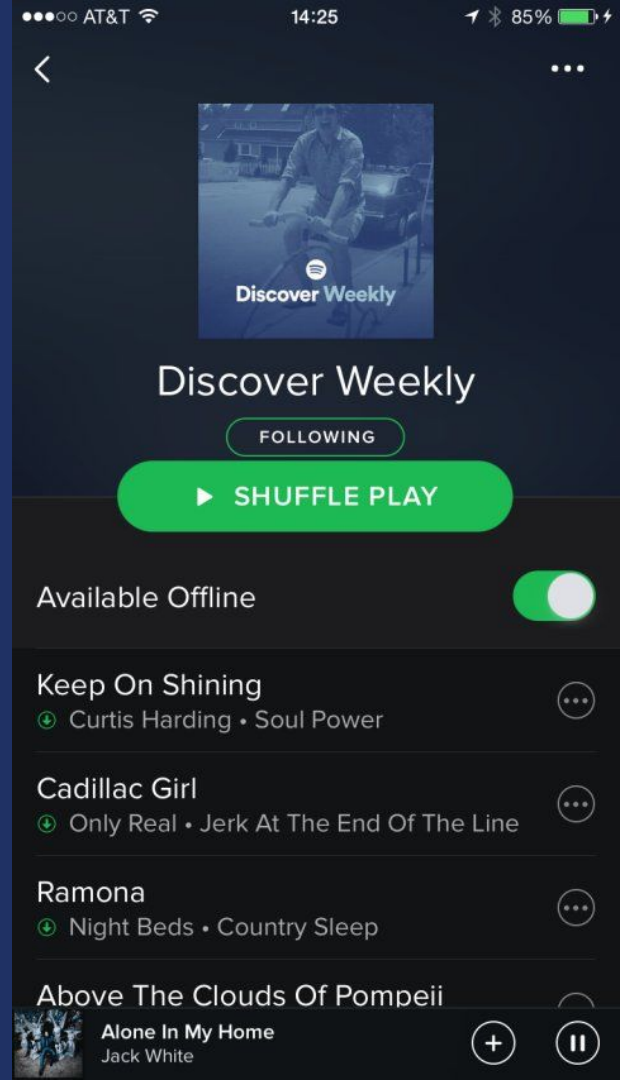
About Us

- 248M+ active users
- 113M+ subscribers
- 50M+ songs
- 500K+ podcasts
- 3B+ playlists



About R&D

- 280+ teams
- 1500+ engineers
- 1200+ microservices
- 8M+ client requests/s
- Half trillion events captured per day



Operations at Spotify



Complete Team
Autonomy



Ops in Teams


Core Infrastructure Org

Golden Path



Centralized Ops

Old Runtime




New Runtime



New Runtime*



Kubernetes: How We're Migrating



Goal/Scope: Run all stateless backend services at Spotify on Kubernetes.
Migrating Spotify's Runtime to Kubernetes: <https://www.youtube.com/watch?v=US4ttqOO-V4>

GA Phase: Golden Path

A The Golden Path for Backend Engineering



Operational Observability and Reliability

Multicloud Strategy: Management

- Automated scripts → Declarative
- Terraform
- How Spotify Accidentally Deleted All its Kube Clusters with No User Impact - David Xia:

<https://kccnceu19.sched.com/event/MQbb/keynote-how-spotify-accidentally-deleted-all-its-kube-clusters-with-no-user-impact-david-xia-infrastructure-engineer-spotify>

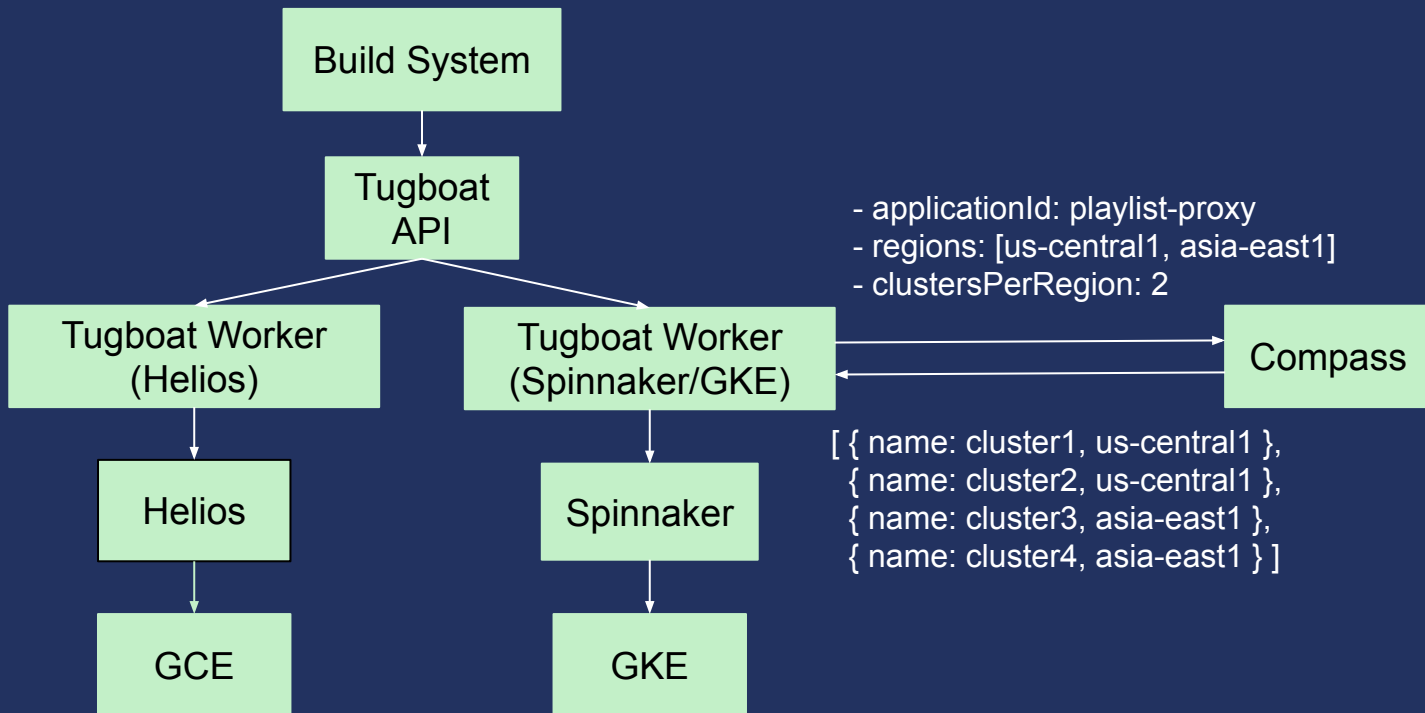


Multicloud Strategy: Management

**Offload cluster
management if possible.**

Multicloud Strategy: Deployment

Tugboat + Compass (deployment layer):



Multicloud Strategy: Compass

- Sticky assignment of apps → clusters
- Main Milestones:
 1. Selection of clusters randomly
 2. Selection of clusters based on available capacity
 3. Ability to deploy an app to N clusters per region

Multicloud Strategy: Deployment

**Keep it simple for
developers to reason about.**

Multicloud: Other Talks

To Infinite Scale and Beyond: Operating Kubernetes Past the Steady State

- Austin Lamon, Spotify & Jago Macleod, Google
- Wednesday (Nov 20) from 3:20pm - 3:55pm

<https://kccncna19.sched.com/event/UaYk/to-infinite-scale-and-beyond-operating-kubernetes-past-the-steady-state-austin-lamon-spotify-jago-macleod-google>

Operational Metrics/Alerts

- Monitoring on integration components
 - error
 - latency
 - saturation
- Smoke test —> periodic deploys
- Metrics for:
 - available cluster capacity
 - status of backups
 - scheduling

Operational Metrics/Alerts

**Establish trust
through monitoring.**

DiRTing

- Disaster Recovery Testing
- Planned and executed DiRTs
 - Cluster deletions
 - Master being down
- Validation and improvement of:
 - Metrics + Alerting
 - Playbooks
 - Backups

DiRTing


**Plan for and
simulate failures.**

Business Impact

Metrics Tracking

- Metrics/Dashboard to track for us:
 - Resource Utilization
 - Cost
 - % apps migrated
 - Developer productivity

Cost Tracking



Business Impact

**Track enough data to
answer questions.**

Self Service

Self-Service Experience: Namespaces

- Checkout the system-resource-manifests monorepo
 - Config that is present on every cluster
- Run a docker image locally with some args
 - Namespace and related resources setup
 - CODEOWNERS setup

```
# generate namespace manifests and supplementary resources
$ docker pull gcr.io/[REDACTED]/gke-onboard
$ docker run -v $(pwd):/generated gcr.io/[REDACTED]/gke-onboard \
  --namespace <namespace> \
  -l <logging_gcp_project> \
  --team <team-name> \
```

Self-Service Experience: Namespaces

**Take complexity for
your developers.**

Self-Service Experience: Apps

```
---  
version: 1  
components:  
- componentId: myservice  
  tugboat:  
    backends:  
      - backend: helios  
      - backend: kubernetes
```

```
version=$(tugboat version create)  
tugboat set-target-version --target-version=$version
```

Self-Service Experience: Apps

**Use existing
tools/abstractions.**

Documentation

- Internal Documentation
- Stackoverflow Enterprise (vast majority of questions on there)
- Backend Golden Path

GKE @ Spotify

Home

Introduction

Getting Started

Roadmap and Use Cases

FAQ

Definitions

Known issues

Feedback

Migration Best Practices

Developers ^

Accessing GCP

Connection Draining

CronJobs

Dealing with Disruptions

Deployment

Troubleshooting Deployments

Event Delivery

ITGC in-scope systems

Java on Kubernetes

Logging

Metrics

Operations

Reference

Scaling


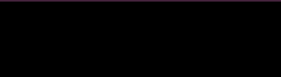
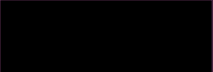
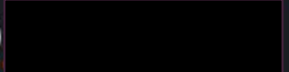
Secrets

Documentation

**You as owners are
not the audience.**

Shared Support

- Developers will always have questions
- US East Coast + Europe
- Shifts covering both
- Ex. Topic for support Slack channel:

For backend services on GKE: <http://docs/gke>. Save time and check  <http://go/so/questions/tagged/gke> | NYC-goalie1: @ 
| NYC-goalie2: @  | STO-goalie: @ 

Shared Support

**Support can be
mutually beneficial.**

Feature Parity

Internal Portal Integration

Visualize Deployment

Currently running on kubernetes in production - Last updated at 1:58 PM

Legend



██████-b04t



Pods: 15

Connect to this cluster



██████-5p4r



Pods: 15



██████-4h0y



Pods: 10



Active Versions

Currently running in production - Last updated at 2:00 PM

Kubernetes

🔍 Search...

↑	REGION	↑	CLUSTER	INSTANCE	COMMIT	STATE	CONSOLE LINK	LOGS
●	██████		██████-4h0y	ranking-86695b7c59-5gqgc	d02eb78f	RUNNING		

Deployment (Canaries)

- Existing canary feature in Tugboat for Helios
- Similar feature on K8s (basically suffix where appropriate with -canary)

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
compass	2	2	2	2	276d
hoth	2	2	2	2	175d
tugboatapi	2	2	2	2	94d
tugboatapi-canary	1	1	1	1	80d
tugboatheliosmonitor	1	1	1	1	272d
tugboatheliosmonitor-canary	1	1	1	1	80d
tugboatheliosworker	3	3	3	3	272d
tugboatheliosworker-canary	1	1	1	1	80d
tugboatnotifications	1	1	1	1	272d
tugboatspinnakerworker	2	2	2	2	354d
tugboatspinnakerworker-canary	1	1	1	1	80d

Feature Parity

Developers need feature parity or better.

ITGC Support

**Delve into the 20 of
the 80/20.**

GA General Requirements

- Operational Solidity
- Business Impact
- Self Service Experience
- Important Feature Parity or better

Future Plans

- Stateful services?
- Data jobs Consolidation
- CronJobs as first class citizens
- Internal Education
- Higher level abstractions

Migration & GA Learnings

- Be mindful and deliberate in dictating terminology.
- Migrate in stages/series of goals that increase in scope.
- Devs will adopt and/or migrate to your Kubernetes offering if it is easy, complete, and compelling.
- Talk to other companies about how they do infrastructure/solve infrastructure problems.



Thanks for listening!

We're hiring!
spotifyjobs.com

James Wen - jameswen@spotify.com
[@rochesterinnyc](#)

KubeCon + CloudNativeCon North America 2019

November 19th, 2019