

The KEP Implementation Journey

Wei Huang, Senior Software Engineer, IBM

KEP Basics

- **KEP: Kubernetes Enhancement Proposal**
 - Consistent feature lifecycle management
 - Serve many audiences
 - <https://github.com/kubernetes/enhancements/tree/master/keps>



kubernetes

CONTRIBUTOR SUMMIT
SAN DIEGO 2019

Before Starting

- **Prerequisite**

- A thoroughly-discussed KEP
- KEP tagged as “implementable”
- Negotiate with SIG(s) leaders about a rough roadmap and align with release schedule
- Reserve your bandwidth

```
---
title: Even Pods Spreading
authors:
  - "@Huang-Wei"
owning-sig: sig-scheduling
reviewers:
  - "@bsalamat"
  - "@lavalamp"
  - "@krmayankk"
approvers:
  - "@bsalamat"
  - "@k82cn"
creation-date: 2019-02-21
last-updated: 2019-05-14
status: implementable
```



Scope, Break Down and Prioritize Items

- **Ask yourself some questions**
 - Is new API needed? If so, what kind of API?
 - In which component(s) to implement? (sometimes the answer is not that obvious, e.g. TaintBasedEviction)
 - Update KEP (esp. API change) if necessary
- 🎉 **Output: an umbrella issue to ensure everyone on the same page**
 - EvenPodsSpread (solo) - [#77284](#)
 - Scheduler Framework (cooperation) - [#83554](#)



API Implementation

- API design pattern


- Top-level API and connect to other API (RuntimeClass, Priority, etc.)

```
// k8s.io/api/core/v1/types.go
type PodSpec struct {
    PriorityClassName string
    RuntimeClassName *string
}
```

— ➔ indirect reference

```
// k8s.io/api/scheduling/v1/types.go
type PriorityClass struct {
    ...
}
```

```
// k8s.io/api/node/v1beta1/types.go
type RuntimeClass struct {
    ...
}
```



kubernetes



CONTRIBUTOR SUMMIT
SAN DIEGO 2019

API Implementation

- API design pattern (cont.)

- Sublevel-API attached to top-level API (Toleration, TopologySpreadConstraint, etc.)

→ direct reference

```
// k8s.io/api/core/v1/types.go
type PodSpec struct {
    Tolerations []Toleration
    // This field is alpha-level and is only honored by clusters that
    // enables the EvenPodsSpread feature.
    TopologySpreadConstraints []TopologySpreadConstraint
}
```

```
// k8s.io/api/core/v1/types.go
type Toleration struct {
    ...
}
```


```
// k8s.io/api/core/v1/types.go
type TopologySpreadConstraint struct {
    ...
}
```


API Implementation

- API design pattern (cont.)
 - CRD (VolumeSnapshot, PodGroup, etc.)

```
// github.com/kubernetes-csi/external-snapshotter/  
// pkg/apis/volumesnapshot/v1beta1/types.go  
type VolumeSnapshot struct {  
    Spec VolumeSnapshotSpec  
}  
  
type VolumeSnapshotSpec struct {  
    Source VolumeSnapshotSource  
}  
  
type VolumeSnapshotSource struct {  
    PersistentVolumeClaimName *string  
}
```

```
// k8s.io/api/core/v1/types.go  
type PersistentVolumeClaimSpec struct {  
    ...  
}
```



kubernetes

CONTRIBUTOR SUMMIT
SAN DIEGO 2019

API Implementation

- API design pattern (cont.)
 - ComponentConfig (Scheduler Framework, etc.)

```
// k8s.io/kube-scheduler/config/v1alpha1/types.go
type KubeSchedulerConfiguration struct {
    ...
    Plugins *Plugins
}

type Plugins struct {
    QueueSort *PluginSet
    Filter *PluginSet
    Score *PluginSet
    Reserve *PluginSet
    Permit *PluginSet
    Bind *PluginSet
    ...
}
```



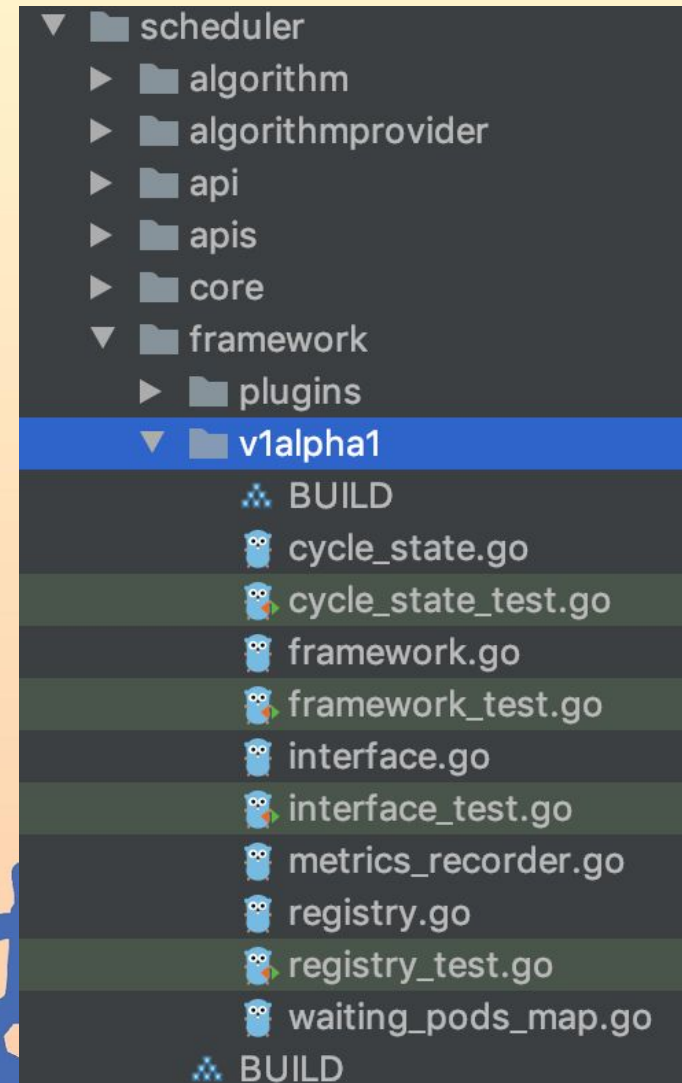
kubernetes



CONTRIBUTOR SUMMIT
SAN DIEGO 2019

API Implementation

- **API design pattern (cont.)**
 - Component internal API
 - Scheduler Framework -
pkg/scheduler/framework/v1alpha1
 - Part of Scheduler Extender -
pkg/scheduler/apis/extender/v1
 - Kubelet
 - pkg/kubelet/apis/podresources/v1alpha1 - protobuf style
 - pkg/kubelet/apis/stats/v1alpha1
 - Usually involves limited or none codegens
 - May not need data persistence



API Implementation

- 🎉 **Output: an API PR ([example](#))**
 - New spec - internal or (and) external
 - GVK (group, version, kind) definition
 - Code Generation
 - Validation, Defaulting
- **Codegen commands (works in 1.16)**

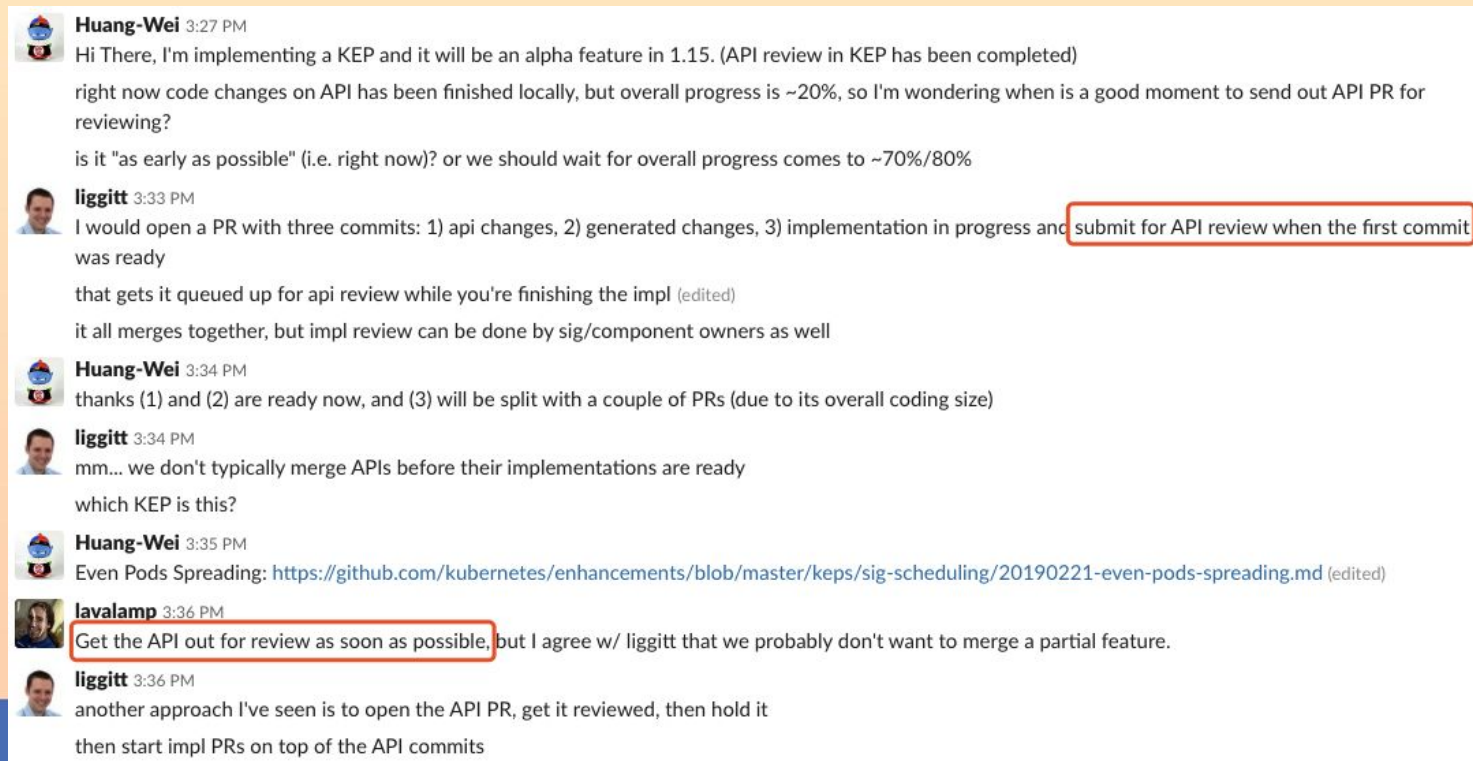
```
make generated_files
hack/update-generated-protobuf.sh
hack/update-codegen.sh
UPDATE_COMPATIBILITY_FIXTURE_DATA=true go test
./vendor/k8s.io/api -run //HEAD &>/dev/null

hack/update-generated-swagger-docs.sh
hack/update-openapi-spec.sh
```

API Implementation

- Tips

- [contributors/devel/sig-architecture/api-conventions.md](https://github.com/kubernetes/contributors/devel/sig-architecture/api-conventions.md)
- [contributors/devel/sig-architecture/api_changes.md](https://github.com/kubernetes/contributors/devel/sig-architecture/api_changes.md)
- Request API code review as early as possible



The screenshot shows a GitHub discussion thread with the following content:

Huang-Wei 3:27 PM
Hi There, I'm implementing a KEP and it will be an alpha feature in 1.15. (API review in KEP has been completed)
right now code changes on API has been finished locally, but overall progress is ~20%, so I'm wondering when is a good moment to send out API PR for reviewing?
is it "as early as possible" (i.e. right now)? or we should wait for overall progress comes to ~70%/80%

liggitt 3:33 PM
I would open a PR with three commits: 1) api changes, 2) generated changes, 3) implementation in progress and submit for API review when the first commit was ready
that gets it queued up for api review while you're finishing the impl (edited)
it all merges together, but impl review can be done by sig/component owners as well

Huang-Wei 3:34 PM
thanks (1) and (2) are ready now, and (3) will be split with a couple of PRs (due to its overall coding size)

liggitt 3:34 PM
mm... we don't typically merge APIs before their implementations are ready
which KEP is this?

Huang-Wei 3:35 PM
Even Pods Spreading: <https://github.com/kubernetes/enhancements/blob/master/keps/sig-scheduling/20190221-even-pods-spreading.md> (edited)

lavalamp 3:36 PM
Get the API out for review as soon as possible, but I agree w/ liggitt that we probably don't want to merge a partial feature.

liggitt 3:36 PM
another approach I've seen is to open the API PR, get it reviewed, then hold it
then start impl PRs on top of the API commits

Core Logic Implementation

- Learn from existing code
- If across multiple components, define the responsibility for each
- Users not using the new feature shouldn't be punished
- Feature gating
 - global feature gate - control core kube components
 - internal feature gating logic - run feature gate check just once

```
// ApplyFeatureGates applies algorithm by feature gates.  
// The returned function is used to restore the state of registered predicates/priorities  
// when this function is called, and should be called in tests which may modify the value  
// of a feature gate temporarily.  
func ApplyFeatureGates() (restore func()) {  
    snapshot := scheduler.RegisteredPredicatesAndPrioritiesSnapshot()  
  
    // Only register EvenPodsSpread predicate & priority if the feature is enabled  
    if utilfeature.DefaultFeatureGate.Enabled(features.EvenPodsSpread) {...}  
  
    // Prioritizes nodes that satisfy pod's resource limits  
    if utilfeature.DefaultFeatureGate.Enabled(features.ResourceLimitsPriorityFunction) {...}  
  
    restore = func() {...}  
    return  
}
```



Core Logic Implementation

- Implement each sub-feature, with tests covered
- Managing PR dependencies and git branches
 - Each PR serves as a standalone functional unit
 - Understand PR dependencies ([#77284](#), [#83554](#)) for better cooperation
 - (optional) Create a separate feature branch on k/k
 - Keep coding, don't stop to await on-going code reviews



Core Logic Implementation

- Managing PR dependencies and git branches (cont.)
 - Develop PRs in parallel to leverage CI
 - **PR1** - commits **A, B**
 - **PR2** - branched off **PR1**, commits **A, B, C**
 - rebase after **PR1** gets merged, so **PR2** only shows commit **C**
 - **PR3** - branched off **PR2**, commits **A, B, C, D, E**
 - rebase after **PR2** gets merged, so **PR3** only shows commits **D, E**
- 🎉 Output: a series of PRs implementing core logic, and sufficient tests
 - Get along well with your friend `git rebase` :)



Test Strategies

- Unit Tests

- **must** be equipped with each PR
 - mandatory for public functions
 - good-to-have for critical private functions
- before `git push`, ensure local UT passes
 - `go test k8s.io/kubernetes/pkg/scheduler/...`
- detect racing condition
 - `go test <pkg> --race`
 - `go test <pkg> --race --count=<number>`
- follow the Golang “[TableDrivenTest](#)” convention
- disable caching
 - `go test <pkg> --count=1`



Test Strategies

- Unit Tests (cont.)

- avoid flakes
 - `go test <pkg> --count=<big number>`
 - be aware of shared variable (featuregate, env variables, tmpdir, etc.)
 - be aware of goroutine leaks
- unable to compose UTs usually implies smelly code
- concrete struct vs. interface
- mocking examples
 - [FakeClientset](#), FakeAPIObj (core API, CRD, ... auto-generated)
 - [FakeClock](#)
 - Test executable - [Real Binary](#) vs. Mock
 - [Test CMD options](#)



Test Strategies

- Unit Tests (cont.)




Dave Cheney
@davecheney

A test is not a unit test if:

- It talks to the database
- It communicates across the network
- It touches the file system
- It can't run at the same time as any of your other unit tests
- You have to do special things to your environment to run it.

—@mfeathers



kubernetes



CONTRIBUTOR SUMMIT
SAN DIEGO 2019

Test Strategies

- **Benchmark Test**

- focus on performance instead of functionalities
- leverage Golang benchmark testing facilities (`*testing.B`)
 - `go test --run=^$ <pkg> --bench .`
(replace `.` with `^<func name>$`)

scheduler: internal data structure optimization #81068

Merged k8s-ci-robot merged 2 commits into `kubernetes:master` from `Huang-Wei:eps-structure-optimize` on Aug 7

Conversation 54 Commits 2 Checks 0 Files changed 8

Huang-Wei commented on Aug 7 • edited ▾ Member + 😊 ...

What type of PR is this?

/kind feature

What this PR does / why we need it:

Optimize the data structure of `topologySpreadMap` `podSpreadCache` to reduce memory and runtime overhead, and also boost the performance of predicates and preemption of `EvenPodsSpread`.

Here is the new structure:

```
// podSpreadCache combines tpKeyToCriticalPaths and tpPairToMatchNum
// to represent:
// (1) critical paths where the least pods are matched on each spread constraint.
// (2) number of pods matched on each spread constraint.
type podSpreadCache struct {
    // We don't record all critical paths here.
    // If there is only one criticalPath, criticalPaths[0] always holds the minimum
    // If there are multiple criticalPaths, keep arbitrary 2 of them holding the min
    tpKeyToCriticalPaths map[string][2]*criticalPath
    // tpPairToMatchNum is keyed with topologyPair, and valued with the number of ma
    tpPairToMatchNum map[topologyPair]int32
}
```

Before

single-constraint-zone-8	100	16717361	ns/op	12650462	B/op	129741	allocs/op
single-constraint-node-8	100	17020951	ns/op	12674237	B/op	132412	allocs/op
two-constraints-zone-node-8	50	22534697	ns/op	16043751	B/op	203792	allocs/op
single-constraint-zone-8	100	17250549	ns/op	12646159	B/op	129737	allocs/op
single-constraint-node-8	100	14650029	ns/op	12672823	B/op	132416	allocs/op
two-constraints-zone-node-8	50	21257976	ns/op	16043452	B/op	203793	allocs/op
single-constraint-zone-8	100	14535822	ns/op	12651489	B/op	129740	allocs/op
single-constraint-node-8	100	14415079	ns/op	12672163	B/op	132414	allocs/op
two-constraints-zone-node-8	100	18504999	ns/op	16041346	B/op	203793	allocs/op

After

single-constraint-zone-8	500	2523942	ns/op	2796400	B/op	60027	allocs/op
single-constraint-node-8	500	2863317	ns/op	2984588	B/op	60051	allocs/op
two-constraints-zone-node-8	300	5202418	ns/op	5753595	B/op	120058	allocs/op
single-constraint-zone-8	500	2574097	ns/op	2795724	B/op	60027	allocs/op
single-constraint-node-8	500	2824376	ns/op	2985523	B/op	60051	allocs/op
two-constraints-zone-node-8	300	4965061	ns/op	5751401	B/op	120059	allocs/op
single-constraint-zone-8	500	2488182	ns/op	2796486	B/op	60027	allocs/op
single-constraint-node-8	500	2922101	ns/op	2988127	B/op	60051	allocs/op
two-constraints-zone-node-8	300	5169896	ns/op	5751460	B/op	120059	allocs/op

Test Strategies

- Benchmark Test - from Unit Test to Benchmark Test

```
func TestFoo(t *testing.T) {
    tests := []struct {
        name     string
        args     ...
        want     retType
    } {
        {
            name: "<description for case1>"
            ...
        }
    }

    for _, tt := range tests {
        t.Run(tt.name, func(t *testing.T) {
            // build dependencies
            r := buildR()
            got := r.Foo(tt.args)
            if !reflect.DeepEqual(got, tt.want) {
                t.Errorf("foo() = %v, want %v", got, tt.want)
            }
        })
    }
}
```

`t *testing.T => b *testing.B`

```
func BenchmarkFoo(b *testing.B) {
    tests := []struct {
        name     string
        args     ...
    } {
        {
            name: "<description for case1>"
            ...
        }
    }

    for _, tt := range tests {
        b.Run(tt.name, func(b *testing.B) {
            // build dependencies
            r := buildR()
            b.ResetTimer()
            for i := 0; i < b.N; i++ {
                // perform the operation we're analyzing
                r.Foo(tt.args)
            }
        })
    }
}
```

remove **want** field

reset timer

`b.N`

No need to verify correctness

Test Strategies

- **Integration Test**

- test sub-component (package) interactions, e.g. EvenPodsSpread
 - [predicates](#), [preemption](#), [priorities](#), [priorities+predicates](#)
- test component interactions (Controller Manager + Scheduler + Kubelet + APIServer), e.g. [TaintBasedEviction](#)
 - node lifecycle manager (part of Controller Manager)
 - scheduler
 - admissions (part of APIServer)
 - some goroutines to simulate kubelet
- usually require etcd/APIServer
- essentially Unit Test



Test Strategies

- Integration Test (cont.)



Bill Kennedy : SEA/-7
@goinggodotnet

Unit Test: Testing a single unit of code. In Go that's a package.

Integration Test: Testing two or more units of code. In Go that's a code path that runs through two or more packages.



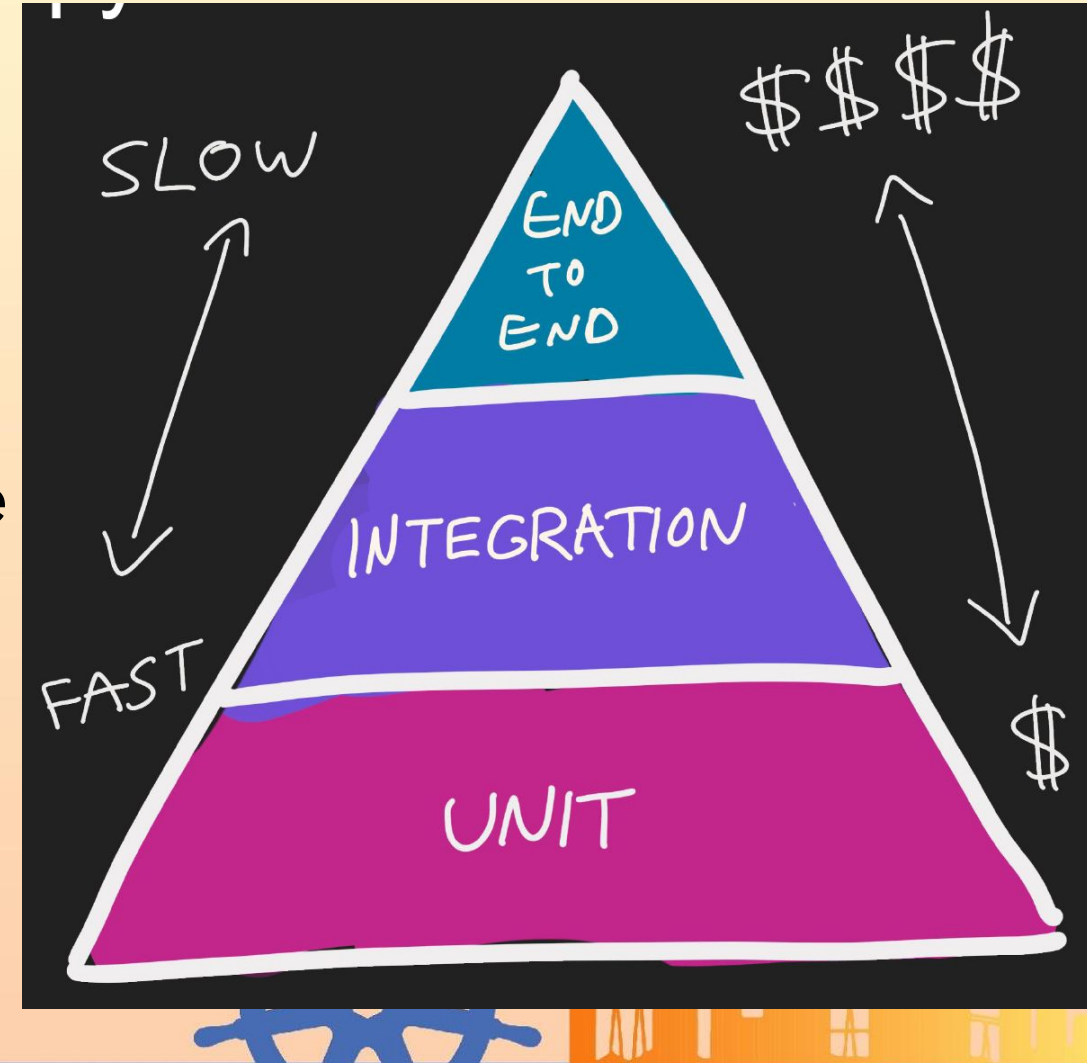
kubernetes

CONTRIBUTOR SUMMIT
SAN DIEGO 2019

Test Strategies

- E2E Test

- uses ginkgo library
- test a real cluster
- black-box testing
- can be promoted to conformance test
- should only test *limited* happy/negative paths of a feature
- optional for alpha feature



Test Strategies


- Enjoy writing tests 😄
 - hard-working and impartial reviewer
 - key factor of an OSS's success



Test Strategies

- **Ad-hoc / Manual tests**

- forget it :)
- how to test a feature manually
 - use kind
 - DIY hacking (only works for control plane components)
 - `START_MODE=nokubelet hack/local-up-cluster.sh`
 - `go build -o /usr/local/bin/hollow-node cmd/kubemark/hollow-node.go`
 - start N hollow-nodes in different terminals
 - `hollow-node --morph=kubelet --name=NodeA --kubeconfig=/var/run/kubernetes/admin.kubeconfig --kubelet-port=10261 --kubelet-read-only-port=10271`



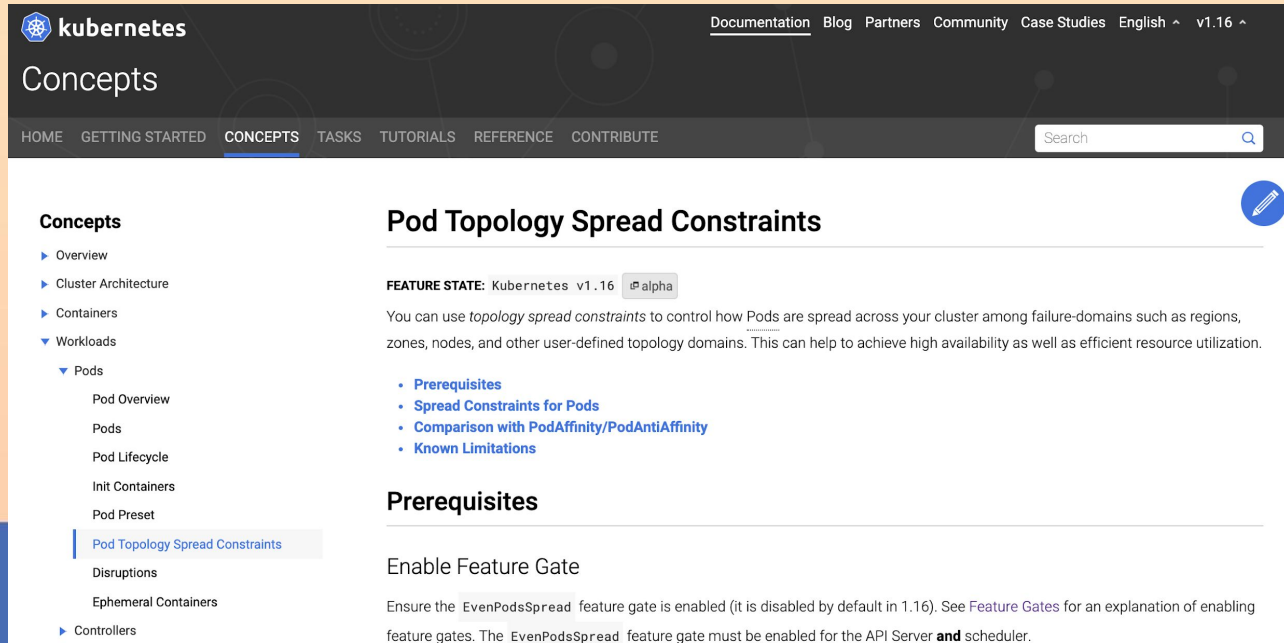
Code Review

- Organize PRs properly
 - explain each section in the PR template, esp. the non-trivial part
 - take **RELEASE NOTES** section seriously! (mistake of EvenPodsSpread)
- Be familiar with Golang [CodeReviewComments](#)
- For a huge PR, resolve comments in incremental commits, and squash them before merge
- Chase reviewers in a reasonable pace



Documents

- API inline documents
 - [example](#) => published at <https://kubernetes.io/docs/reference/generated/kubernetes-api/v1.16>
- kubernetes/website documents
 - [example](#) => published at <https://kubernetes.io>



Tools/Commands

- **text searching tool:** `ag` (a.k.a the silver searcher)
- **git tools:** `tig`, `grv`
- **git commands and aliases** (some come up with oh-my-zsh)
 - switch to latest-used branch: `git checkout -`
 - squash/amend commits: `git rebase -i`
 - ban ~~`git merge`~~ as the “Merge...” commit pollutes the checkin history
 - searching/identifying: `git blame`, `git bisect`
 - aliases: `gppush`, `glum`, `gbda`, `gco -b`, `gcp`, etc.
- **IDE**
 - Goland vs. VSCode



kubernetes



CONTRIBUTOR SUMMIT
SAN DIEGO 2019



Thanks!

Github: @Huang-Wei

Slack: @Huang-Wei

Twitter: @hweicdl