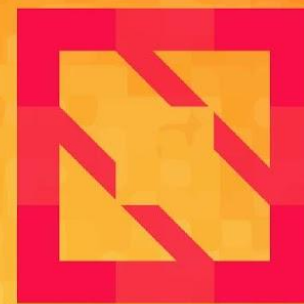




KubeCon



CloudNativeCon

North America 2019





KubeCon



CloudNativeCon

North America 2019

Weighing a Cloud:

Measuring Your Kubernetes Clusters

Han Kang, Google & Elana Hashman, Red Hat



Who are we?



Han Kang

Senior Software Engineer

- Cluster Ops Lead at Google
- SIG API-Machinery and SIG Instrumentation Member
- Twitter: **@LogicalHan**
- GitHub: **@logicalhan**



Elana Hashman

Principal Site Reliability Engineer

- Tech Lead on Azure Red Hat OpenShift Team
- SIG Instrumentation Member
- Twitter: **@ehashdn**
- GitHub: **@ehashman**

What we are going to cover



KubeCon



CloudNativeCon

North America 2019

- How instrumentation works in Kubernetes
- Kubernetes control plane instrumentation
- Real-world debugging!
- Metric usability and SIG Instrumentation roadmap



KubeCon



CloudNativeCon


North America 2019

How Kubernetes Instrumentation Works



Prometheus

Kubernetes components integrate with Prometheus, a time-series based monitoring and alerting toolkit.



Prometheus Data Model

Timeseries

`up{job="kube-apiserver",instance="api-1"}`

Value

1

Types of metric values:

- Counters
- Gauges
- Summaries
- Histograms

Dimensions of Measurement

1. Availability

- `up{job="kubernetes-apiservers"}`

2. Latency

- `apiserver_request_latency_seconds`

3. Capacity

- `apiserver_request_total`

4. Errors

- `apiserver_dropped_requests_total`

Using Prometheus Metrics

Prometheus query language (PromQL) powers metrics analysis and aggregation

- **For prototyping and exploration:** use the Prometheus UI
- **For permanent dashboards:** attach a Prometheus data source to Grafana
- **For alerting:** set up the Prometheus Alert Manager
- **For arbitrary queries and processing:** query the Prometheus API

Differential Diagnoses

- Lots of very different issues might manifest the same way
 - e.g. “a node is offline” -- but why?
- A single symptom is not sufficient to form a diagnosis
- Metrics can show **how** something is failing, but not **why**
- We must track down root causes with multiple data sources

Full-Stack Debugging



KubeCon



CloudNativeCon

North America 2019

- Metrics can guide you to what you should look at next
- Not just metrics!
 - log files
 - audit logs
 - events
 - etcd (cluster database) dumps
- Metrics are most effective when you understand the context in which they were produced.




KubeCon



CloudNativeCon

North America 2019






KubeCon



CloudNativeCon

North America 2019





KubeCon



CloudNativeCon

North America 2019

Kubernetes Control Plane Instrumentation



Control Plane




KubeCon



CloudNativeCon

North America 2019



Kubelet




KubeCon



CloudNativeCon

North America 2019



LIVENESS PROBE

Master Kubelet




KubeCon



CloudNativeCon

North America 2019



Introspecting Components



KubeCon



CloudNativeCon

North America 2019

1. health check endpoint(s)
2. metrics
3. logs



Introspecting Components




KubeCon



CloudNativeCon

North America 2019

```
$ curl localhost:10251/healthz?verbose  
[+]leaderElection ok  
healthz check passed
```



KAS (Kube-apiserver)




KubeCon




CloudNativeCon

North America 2019




Kube-apiserver



- `kubectl <command> -v=9`
...
`round_tripper.go:386] curl <some headers>`
`'https://masterip/api/v1/components tatuses?limit=500'`

Kube-apiserver



- `kubectl <command> -v=9`
- `kube-apiserver.log`
- `/metrics`
- `health endpoints`
 - `localhost:8080/healthz?verbose`
 - `localhost:8080/livez (v1.16+)`
 - `localhost:8080/readyz (v1.16+)`
- `audit-logs`

Etcd

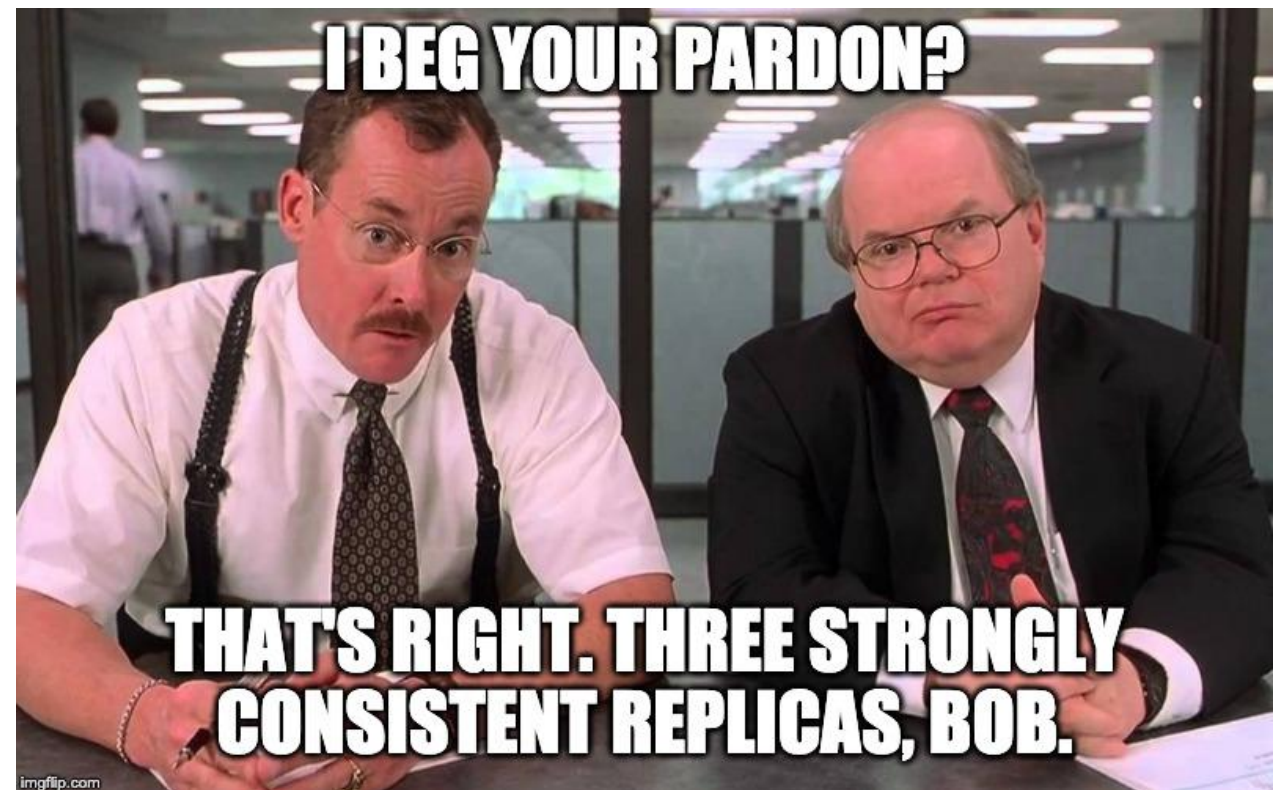



KubeCon




CloudNativeCon

North America 2019



Etcd



- `etcdctl`
- `auger`
- `/metrics`
- `/health`
- `etcd.log`



KubeCon



CloudNativeCon

North America 2019

Real-world Debugging



Kubelet Example



KubeCon



CloudNativeCon

North America 2019

Problem:
Node is down



Kubelet Example

- Obvious: Prometheus scrape job is down
`up{job="kube-nodes"} != 1`
- Less obvious: Grey failure indicated by unusually slow scrape time
`scrape_duration_seconds{job="kube-nodes"} > 2`

Kube-apiserver Example

Problem:
Crash-looping
kube-apiserver



Kube-apiserver Example

Detection Strategies:

1. Directly monitor kube-apiserver health endpoints
2. Alerting based off master kubelets 'metrics/probes'

Kube-apiserver Example

output of kubelet's metrics/probes

HELP prober_probe_total Cumulative number of a liveness or readiness probe for a container by result.

TYPE prober_probe_total counter

prober_probe_total{container="kube-apiserver",probe_type="Liveness",result="failed"} 10

prober_probe_total{container="kube-apiserver",probe_type="Liveness",result="successful"} 26457

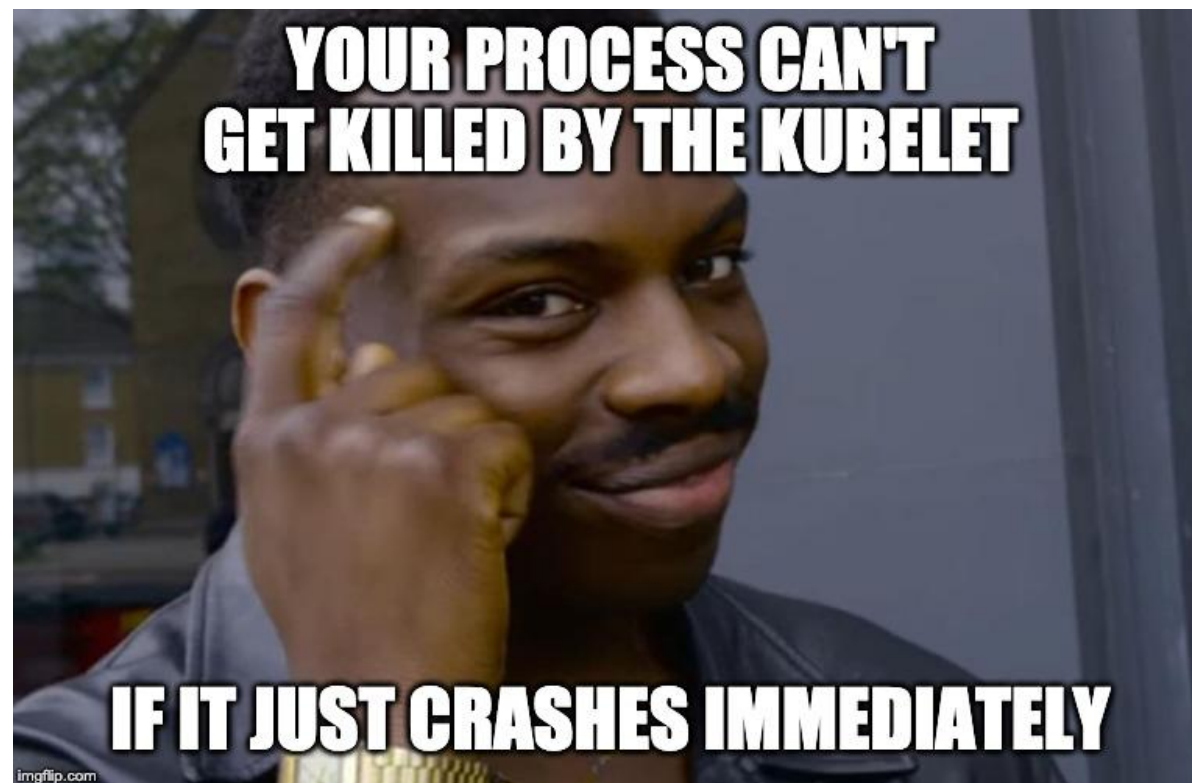
prober_probe_total{container="kube-apiserver",probe_type="Readiness",result="failed"} 16

prober_probe_total{container="kube-apiserver",probe_type="Readiness",result="successful"} 26458

Kube-apiserver Example

Possible reasons:

- a. kubelet in repair mode
- b. kubelet initiated crashloops



Kube-apiserver Example



KubeCon



CloudNativeCon

North America 2019

kube-apiserver /healthz

```
$ curl localhost:8080/healthz?verbose
```

```
[+]ping ok
```


```
[+]log ok
```

```
[-]etcd failed: reason withheld
```

```
..... ok
```

```
[+]autoregister-completion ok
```

```
healthz check failed
```



Etcd Example

Detection Strategies:

- ~~1. Directly monitor etcd health endpoints~~
- ~~2. Directly monitor kube apiserver health endpoints~~
3. Alerting based off master kubelets 'metrics/probes'

Etcd Example

```
# HELP etcd_object_counts Number of stored objects at the time of last check split by kind.  
# TYPE etcd_object_counts gauge  
etcd_object_counts{resource="somecrd"} 1000000
```

Storage size limit

(<https://github.com/etcd-io/etcd/blob/release-3.4/Documentation/dev-guide/limit.md>)

The default storage size limit is 2GB, configurable with `--quota-backend-bytes` flag. 8GB is a suggested maximum size for normal environments and etcd warns at startup if the configured value exceeds it.

Etcd Example



KubeCon



CloudNativeCon

North America 2019

```
etcd_object_counts{resource="somecrd"} 1  
apiserver_request_count{resource="somecrd", verb="UPDATE"} 1200
```

Etcd Example

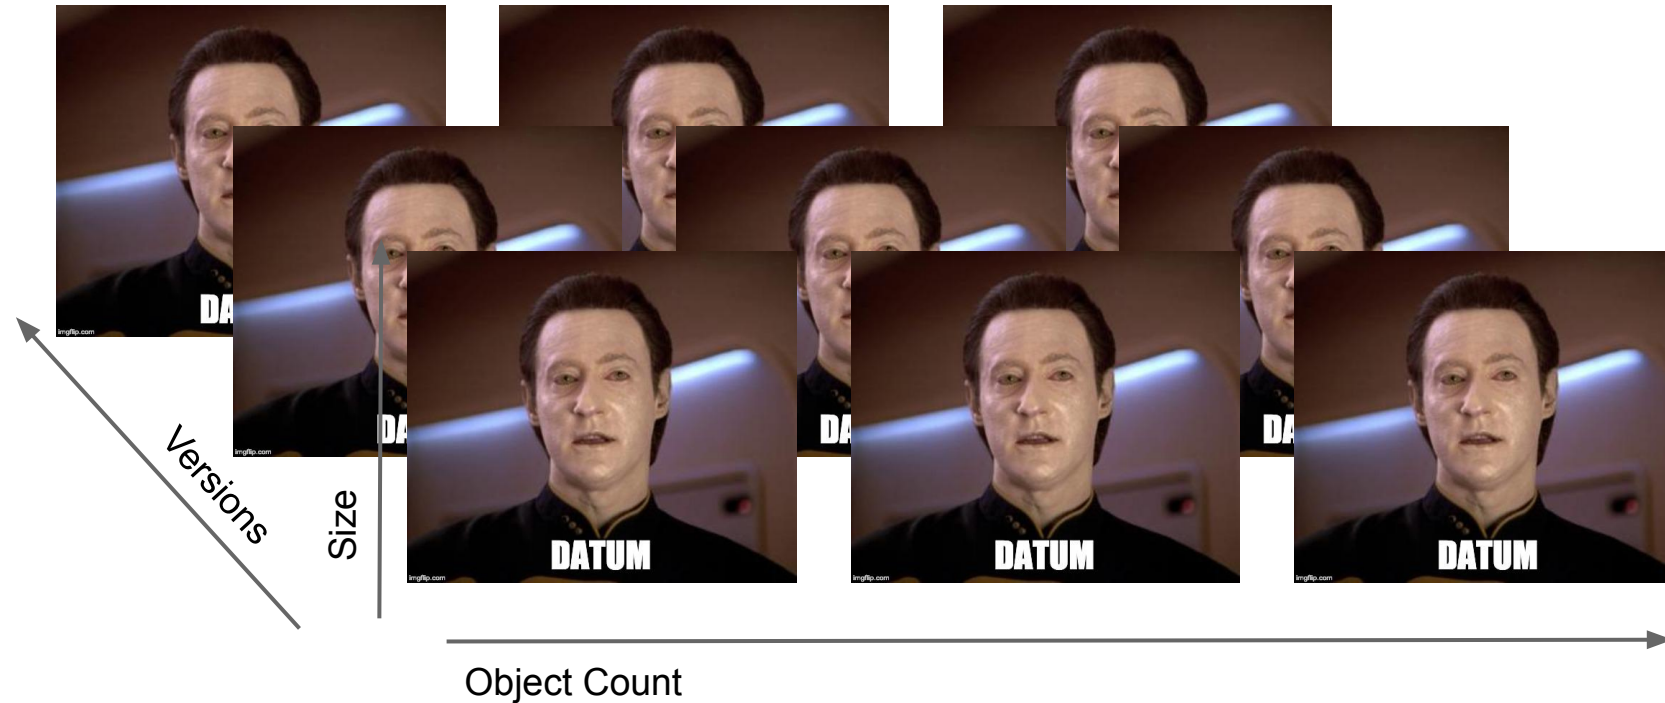


KubeCon



CloudNativeCon

North America 2019



```
# (Revisited):
```

```
etcd_object_counts{resource="somecrd.io"} 1
```

```
apiserver_request_count{resource="somecrd.io", verb="UPDATE"}
```

```
1200
```

Etcd Example



KubeCon



CloudNativeCon

North America 2019

```
$ kubectl get -ojson somecrd.io datum | wc -c
```

```
$ auger extract -f <dbfile> -k <key> | wc -c
```

Kube-apiserver Example



KubeCon



CloudNativeCon

North America 2019

Problem:
API-servers are
slow.



Another kube-apiserver example

Problem: API servers are slow

- Obvious: p99 request latency is high

```
histogram_quantile(  
    0.99,  
    sum(rate(apiserver_request_latencies_bucket[1m]))  
        by (le, verb)  
)
```

Another kube-apiserver example




KubeCon



CloudNativeCon

North America 2019



Another kube-apiserver example

Problem: API servers are slow

- Less obvious: API server metrics prior to 1.14 release are limited to buckets between 125ms and 8s

Another kube-apiserver example



KubeCon



CloudNativeCon

North America 2019

Adjust buckets in apiserver request latency metrics

[Browse files](#)

master (#73638) v1.15.0-alpha.0 ... v1.14.0-alpha.3



wojtek-t committed on Feb 1

1 parent [a3c14ec](#)

commit d0508c7e872f60826d68c58c458cfd865554b486

Showing **1 changed file** with **5 additions** and **2 deletions**.

Unified

Split

7 staging/src/k8s.io/apiserver/pkg/endpoints/metrics/metrics.go

[View file](#)



```
8 +72,11 @@ var (  
    prometheus.HistogramOpts{  
        Name: "apiserver_request_latency_seconds",  
        Help: "Response latency distribution in seconds for each verb, group, version, resource, subresource, scope and component",  
        // Use buckets ranging from 125 ms to 8 seconds.  
        Buckets: prometheus.ExponentialBuckets(0.125, 2.0, 7),  
        // This metric is used for verifying api call latencies SLO,  
        // as well as tracking regressions in this aspects.  
        // Thus we customize buckets significantly, to empower both usecases.  
        Buckets: []float64{0.05, 0.1, 0.15, 0.2, 0.25, 0.3, 0.35, 0.4, 0.45, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0,  
                        1.25, 1.5, 1.75, 2.0, 2.5, 3.0, 3.5, 4.0, 4.5, 5, 6, 7, 8, 9, 10, 15, 20, 25, 30, 40, 50, 60},  
    },  
    []string{"verb", "group", "version", "resource", "subresource", "scope", "component"},  
)
```




KubeCon



CloudNativeCon

North America 2019

Metric Usability & SIG Instrumentation



Handling metric issues



KubeCon



CloudNativeCon

North America 2019

- SIG Instrumentation needs to be able to fix metric bugs and issues
- Updating metrics between releases could break monitoring stacks
- Bad metrics can't be disabled, requiring a full upgrade to address
- How can we coordinate developers to address this and responsibly communicate to end users?

Metrics Overhaul (1.14)

- Many **broken metrics** were identified
 - Labels did not match instrumentation guidelines, couldn't be joined
 - Wrong data types prevented aggregation
 - Units were not standardized
- **Fixes rolled out** in the 1.14 release
- SIG Instrumentation KEP: “Kubernetes Metrics Overhaul”

Metric Stability Framework

- SIG Instrumentation KEP: “Kubernetes Control Plane Metrics Stability”
- **Treat metrics as a proper API:** multi-release notice period for changes to stable metrics
- **Deprecation lifecycle:** slowly phase out obsolete metrics across releases before deletion
- **Enforcing Stability:** metrics migration, static analysis for stability validation, beta enforcements

Stability Metadata



KubeCon



CloudNativeCon

North America 2019

```
var rpcDurations = metrics.NewSummary(  
    metrics.SummaryOpts{  
        Name:      "rpc_durations_seconds",  
        Help:      "RPC latency distributions.",  
        StabilityLevel: metrics.STABLE,  
        DeprecatedVersion: "1.15",  
    },  
)  
)
```

More to come



KubeCon



CloudNativeCon

North America 2019

- Stable metric criteria and promotion
- Distributed tracing
- Structured logs
- More metric improvements!



KubeCon



CloudNativeCon

North America 2019

Questions?





KubeCon



CloudNativeCon

North America 2019



Image Citations

- Slide 16 : Title: *Liveness Probe Meme*; Site: Meme Generator; URL: <https://imgflip.com/memegenerator>; Date: 11/15/19; Publisher: imgflip
- Slide 18 : Title: *Count on me Meme*; Site: Meme Generator; URL: <https://imgflip.com/memegenerator>; Date: 11/15/19; Publisher: imgflip
- Slide 19 : Title: *Only one health check Meme*; Site: Meme Generator; URL: <https://imgflip.com/memegenerator>; Date: 11/15/19; Publisher: imgflip
- Slide 20 : Title: *Talk to the hand Meme*; Site: Meme Generator; URL: <https://imgflip.com/memegenerator>; Date: 11/15/19; Publisher: imgflip
- Slide 23 : Title: *Etcd Meme*; Site: Meme Generator; URL: <https://imgflip.com/memegenerator>; Date: 11/15/19; Publisher: imgflip
- Slide 26 : Title: *Bambi Meme*; Site: Meme Generator; URL: <https://imgflip.com/memegenerator>; Date: 11/15/19; Publisher: imgflip
- Slide 28: Title: *Oh nos Meme*; Site: Meme Generator; URL: <https://imgflip.com/memegenerator>; Date: 11/15/19; Publisher: imgflip
- Slide 31 : Title: *Can't crash a crashed process Meme*; Site: Meme Generator; URL: <https://imgflip.com/memegenerator>; Date: 11/15/19; Publisher: imgflip
- Slide 32: Title: *Causes other crashloops Meme*; Site: Meme Generator; URL: <https://imgflip.com/memegenerator>; Date: 11/15/19; Publisher: imgflip
- Slide 36 : Title: *Datum Meme*; Site: Meme Generator; URL: <https://imgflip.com/memegenerator>; Date: 11/15/19; Publisher: imgflip
- Slide 38 : Title: *Silly latency metric Meme*; Site: Meme Generator; URL: <https://imgflip.com/memegenerator>; Date: 11/15/19; Publisher: imgflip